

# T81-558: Applications of Deep Neural Networks

## Module 4: Training for Tabular Data

- Instructor: [Jeff Heaton](#), McKelvey School of Engineering, [Washington University in St. Louis](#)
- For more information visit the [class website](#).

## Module 4 Material

- Part 4.1: Encoding a Feature Vector for Keras Deep Learning [\[Video\]](#) [\[Notebook\]](#)
- Part 4.2: Keras Multiclass Classification for Deep Neural Networks with ROC and AUC [\[Video\]](#) [\[Notebook\]](#)
- **Part 4.3: Keras Regression for Deep Neural Networks with RMSE** [\[Video\]](#) [\[Notebook\]](#)
- Part 4.4: Backpropagation, Nesterov Momentum, and ADAM Neural Network Training [\[Video\]](#) [\[Notebook\]](#)
- Part 4.5: Neural Network RMSE and Log Loss Error Calculation from Scratch [\[Video\]](#) [\[Notebook\]](#)

## Google CoLab Instructions

The following code ensures that Google CoLab is running the correct version of TensorFlow.

```
In [1]: try:
        %tensorflow_version 2.x
        COLAB = True
        print("Note: using Google CoLab")
    except:
        print("Note: not using Google CoLab")
        COLAB = False
```

Note: not using Google CoLab

## Part 4.3: Keras Regression for Deep Neural Networks with RMSE

We evaluate regression results differently than classification. Consider the following code that trains a neural network for regression on the data set **jh-simple-dataset.csv**. We begin by preparing the data set.

```
In [2]: import pandas as pd
from scipy.stats import zscore
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# Read the data set
df = pd.read_csv(
    "https://data.heatonresearch.com/data/t81-558/jh-simple-dataset.csv",
    na_values=['NA', '?'])

# Generate dummies for job
df = pd.concat([df, pd.get_dummies(df['job'], prefix="job")], axis=1)
df.drop('job', axis=1, inplace=True)

# Generate dummies for area
df = pd.concat([df, pd.get_dummies(df['area'], prefix="area")], axis=1)
df.drop('area', axis=1, inplace=True)

# Generate dummies for product
df = pd.concat([df, pd.get_dummies(df['product'], prefix="product")], axis=1)
df.drop('product', axis=1, inplace=True)

# Missing values for income
med = df['income'].median()
df['income'] = df['income'].fillna(med)

# Standardize ranges
df['income'] = zscore(df['income'])
df['aspect'] = zscore(df['aspect'])
df['save_rate'] = zscore(df['save_rate'])
df['subscriptions'] = zscore(df['subscriptions'])

# Convert to numpy - Classification
x_columns = df.columns.drop('age').drop('id')
x = df[x_columns].values
y = df['age'].values

# Create train/test
x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.25, random_state=42)
```

Next, we create a neural network to fit the data we just loaded.

```
In [3]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Activation
        from tensorflow.keras.callbacks import EarlyStopping

        # Build the neural network
        model = Sequential()
        model.add(Dense(25, input_dim=x.shape[1], activation='relu')) # Hidden 1
        model.add(Dense(10, activation='relu')) # Hidden 2
        model.add(Dense(1)) # Output
        model.compile(loss='mean_squared_error', optimizer='adam')
        monitor = EarlyStopping(monitor='val_loss', min_delta=1e-3,
                                patience=5, verbose=1, mode='auto',
                                restore_best_weights=True)
        model.fit(x_train,y_train,validation_data=(x_test,y_test),
                  callbacks=[monitor],verbose=2,epochs=1000)
```

Train on 1500 samples, validate on 500 samples

Epoch 1/1000	1500/1500	- 1s	- loss: 1905.4454	- val_loss: 1628.1341
Epoch 2/1000	1500/1500	- 0s	- loss: 1331.4213	- val_loss: 889.0575
Epoch 3/1000	1500/1500	- 0s	- loss: 554.8426	- val_loss: 303.7261
Epoch 4/1000	1500/1500	- 0s	- loss: 276.2087	- val_loss: 241.2495
Epoch 5/1000	1500/1500	- 0s	- loss: 232.2832	- val_loss: 208.2143
Epoch 6/1000	1500/1500	- 0s	- loss: 198.5331	- val_loss: 179.5262
Epoch 7/1000	1500/1500	- 0s	- loss: 169.0791	- val_loss: 154.5270
Epoch 8/1000	1500/1500	- 0s	- loss: 144.1286	- val_loss: 132.8691
Epoch 9/1000	1500/1500	- 0s	- loss: 122.9873	- val_loss: 115.0928
Epoch 10/1000	1500/1500	- 0s	- loss: 104.7249	- val_loss: 98.7375
Epoch 11/1000	1500/1500	- 0s	- loss: 89.8292	- val_loss: 86.2749
Epoch 12/1000	1500/1500	- 0s	- loss: 77.3071	- val_loss: 75.0022
Epoch 13/1000	1500/1500	- 0s	- loss: 67.0604	- val_loss: 66.1396
Epoch 14/1000	1500/1500	- 0s	- loss: 58.9584	- val_loss: 58.4367
Epoch 15/1000	1500/1500	- 0s	- loss: 51.2491	- val_loss: 52.7136
Epoch 16/1000	1500/1500	- 0s	- loss: 45.1765	- val_loss: 46.5179
Epoch 17/1000	1500/1500	- 0s	- loss: 39.8843	- val_loss: 41.3721
Epoch 18/1000	1500/1500	- 0s	- loss: 35.1468	- val_loss: 37.2132
Epoch 19/1000	1500/1500	- 0s	- loss: 31.1755	- val_loss: 33.0697
Epoch 20/1000	1500/1500	- 0s	- loss: 27.6307	- val_loss: 30.3131
Epoch 21/1000	1500/1500	- 0s	- loss: 24.8457	- val_loss: 26.9474
Epoch 22/1000	1500/1500	- 0s	- loss: 22.4056	- val_loss: 24.3656
Epoch 23/1000	1500/1500	- 0s	- loss: 20.3071	- val_loss: 22.1642
Epoch 24/1000	1500/1500	- 0s	- loss: 18.5446	- val_loss: 20.4782
Epoch 25/1000	1500/1500	- 0s	- loss: 17.1571	- val_loss: 18.8670
Epoch 26/1000	1500/1500	- 0s	- loss: 15.9407	- val_loss: 17.6862
Epoch 27/1000	1500/1500	- 0s	- loss: 14.9866	- val_loss: 16.5275
Epoch 28/1000				

1500/1500 - 0s - loss: 14.1251 - val\_loss: 15.6342  
Epoch 29/1000  
1500/1500 - 0s - loss: 13.4655 - val\_loss: 14.8625  
Epoch 30/1000  
1500/1500 - 0s - loss: 12.8994 - val\_loss: 14.2826  
Epoch 31/1000  
1500/1500 - 0s - loss: 12.5566 - val\_loss: 13.6121  
Epoch 32/1000  
1500/1500 - 0s - loss: 12.0077 - val\_loss: 13.3087  
Epoch 33/1000  
1500/1500 - 0s - loss: 11.5357 - val\_loss: 12.6593  
Epoch 34/1000  
1500/1500 - 0s - loss: 11.2365 - val\_loss: 12.1849  
Epoch 35/1000  
1500/1500 - 0s - loss: 10.8074 - val\_loss: 11.9388  
Epoch 36/1000  
1500/1500 - 0s - loss: 10.5593 - val\_loss: 11.4006  
Epoch 37/1000  
1500/1500 - 0s - loss: 10.2093 - val\_loss: 10.9751  
Epoch 38/1000  
1500/1500 - 0s - loss: 9.8386 - val\_loss: 10.8651  
Epoch 39/1000  
1500/1500 - 0s - loss: 9.5938 - val\_loss: 10.5728  
Epoch 40/1000  
1500/1500 - 0s - loss: 9.1488 - val\_loss: 9.8661  
Epoch 41/1000  
1500/1500 - 0s - loss: 8.8920 - val\_loss: 9.5228  
Epoch 42/1000  
1500/1500 - 0s - loss: 8.5156 - val\_loss: 9.1506  
Epoch 43/1000  
1500/1500 - 0s - loss: 8.2628 - val\_loss: 8.9486  
Epoch 44/1000  
1500/1500 - 0s - loss: 7.9219 - val\_loss: 8.5034  
Epoch 45/1000  
1500/1500 - 0s - loss: 7.7077 - val\_loss: 8.0760  
Epoch 46/1000  
1500/1500 - 0s - loss: 7.3165 - val\_loss: 7.6620  
Epoch 47/1000  
1500/1500 - 0s - loss: 7.0259 - val\_loss: 7.4933  
Epoch 48/1000  
1500/1500 - 0s - loss: 6.7422 - val\_loss: 7.0583  
Epoch 49/1000  
1500/1500 - 0s - loss: 6.5163 - val\_loss: 6.8024  
Epoch 50/1000  
1500/1500 - 0s - loss: 6.2633 - val\_loss: 7.3045  
Epoch 51/1000  
1500/1500 - 0s - loss: 6.0029 - val\_loss: 6.2712  
Epoch 52/1000  
1500/1500 - 0s - loss: 5.6791 - val\_loss: 5.9342  
Epoch 53/1000  
1500/1500 - 0s - loss: 5.4798 - val\_loss: 6.0110  
Epoch 54/1000  
1500/1500 - 0s - loss: 5.2115 - val\_loss: 5.3928  
Epoch 55/1000  
1500/1500 - 0s - loss: 4.9592 - val\_loss: 5.2215  
Epoch 56/1000

1500/1500 - 0s - loss: 4.7189 - val\_loss: 5.0103  
Epoch 57/1000  
1500/1500 - 0s - loss: 4.4683 - val\_loss: 4.7098  
Epoch 58/1000  
1500/1500 - 0s - loss: 4.2650 - val\_loss: 4.5259  
Epoch 59/1000  
1500/1500 - 0s - loss: 4.0953 - val\_loss: 4.4263  
Epoch 60/1000  
1500/1500 - 0s - loss: 3.8027 - val\_loss: 4.1103  
Epoch 61/1000  
1500/1500 - 0s - loss: 3.5759 - val\_loss: 3.7770  
Epoch 62/1000  
1500/1500 - 0s - loss: 3.3755 - val\_loss: 3.5737  
Epoch 63/1000  
1500/1500 - 0s - loss: 3.1781 - val\_loss: 3.4833  
Epoch 64/1000  
1500/1500 - 0s - loss: 3.0001 - val\_loss: 3.2246  
Epoch 65/1000  
1500/1500 - 0s - loss: 2.7691 - val\_loss: 3.1021  
Epoch 66/1000  
1500/1500 - 0s - loss: 2.6227 - val\_loss: 2.8215  
Epoch 67/1000  
1500/1500 - 0s - loss: 2.4682 - val\_loss: 2.7528  
Epoch 68/1000  
1500/1500 - 0s - loss: 2.3243 - val\_loss: 2.5394  
Epoch 69/1000  
1500/1500 - 0s - loss: 2.1664 - val\_loss: 2.3886  
Epoch 70/1000  
1500/1500 - 0s - loss: 2.0377 - val\_loss: 2.2536  
Epoch 71/1000  
1500/1500 - 0s - loss: 1.8845 - val\_loss: 2.2354  
Epoch 72/1000  
1500/1500 - 0s - loss: 1.7931 - val\_loss: 2.0831  
Epoch 73/1000  
1500/1500 - 0s - loss: 1.6889 - val\_loss: 1.8866  
Epoch 74/1000  
1500/1500 - 0s - loss: 1.5820 - val\_loss: 1.7964  
Epoch 75/1000  
1500/1500 - 0s - loss: 1.5085 - val\_loss: 1.7138  
Epoch 76/1000  
1500/1500 - 0s - loss: 1.4159 - val\_loss: 1.6468  
Epoch 77/1000  
1500/1500 - 0s - loss: 1.3606 - val\_loss: 1.5906  
Epoch 78/1000  
1500/1500 - 0s - loss: 1.2652 - val\_loss: 1.5063  
Epoch 79/1000  
1500/1500 - 0s - loss: 1.1937 - val\_loss: 1.4506  
Epoch 80/1000  
1500/1500 - 0s - loss: 1.1180 - val\_loss: 1.4817  
Epoch 81/1000  
1500/1500 - 0s - loss: 1.1412 - val\_loss: 1.2800  
Epoch 82/1000  
1500/1500 - 0s - loss: 1.0385 - val\_loss: 1.2412  
Epoch 83/1000  
1500/1500 - 0s - loss: 0.9846 - val\_loss: 1.1891  
Epoch 84/1000

1500/1500 - 0s - loss: 0.9937 - val\_loss: 1.1322  
Epoch 85/1000  
1500/1500 - 0s - loss: 0.8915 - val\_loss: 1.0847  
Epoch 86/1000  
1500/1500 - 0s - loss: 0.8562 - val\_loss: 1.1110  
Epoch 87/1000  
1500/1500 - 0s - loss: 0.8468 - val\_loss: 1.0686  
Epoch 88/1000  
1500/1500 - 0s - loss: 0.7947 - val\_loss: 0.9805  
Epoch 89/1000  
1500/1500 - 0s - loss: 0.7807 - val\_loss: 0.9463  
Epoch 90/1000  
1500/1500 - 0s - loss: 0.7502 - val\_loss: 0.9965  
Epoch 91/1000  
1500/1500 - 0s - loss: 0.7529 - val\_loss: 0.9532  
Epoch 92/1000  
1500/1500 - 0s - loss: 0.6857 - val\_loss: 0.8712  
Epoch 93/1000  
1500/1500 - 0s - loss: 0.6717 - val\_loss: 0.8498  
Epoch 94/1000  
1500/1500 - 0s - loss: 0.6869 - val\_loss: 0.8518  
Epoch 95/1000  
1500/1500 - 0s - loss: 0.6626 - val\_loss: 0.8275  
Epoch 96/1000  
1500/1500 - 0s - loss: 0.6308 - val\_loss: 0.7850  
Epoch 97/1000  
1500/1500 - 0s - loss: 0.6056 - val\_loss: 0.7708  
Epoch 98/1000  
1500/1500 - 0s - loss: 0.5991 - val\_loss: 0.7643  
Epoch 99/1000  
1500/1500 - 0s - loss: 0.6102 - val\_loss: 0.8104  
Epoch 100/1000  
1500/1500 - 0s - loss: 0.5647 - val\_loss: 0.7227  
Epoch 101/1000  
1500/1500 - 0s - loss: 0.5474 - val\_loss: 0.7107  
Epoch 102/1000  
1500/1500 - 0s - loss: 0.5395 - val\_loss: 0.6847  
Epoch 103/1000  
1500/1500 - 0s - loss: 0.5350 - val\_loss: 0.7383  
Epoch 104/1000  
1500/1500 - 0s - loss: 0.5551 - val\_loss: 0.6698  
Epoch 105/1000  
1500/1500 - 0s - loss: 0.5032 - val\_loss: 0.6520  
Epoch 106/1000  
1500/1500 - 0s - loss: 0.5418 - val\_loss: 0.7518  
Epoch 107/1000  
1500/1500 - 0s - loss: 0.4949 - val\_loss: 0.6307  
Epoch 108/1000  
1500/1500 - 0s - loss: 0.5166 - val\_loss: 0.6741  
Epoch 109/1000  
1500/1500 - 0s - loss: 0.4992 - val\_loss: 0.6195  
Epoch 110/1000  
1500/1500 - 0s - loss: 0.4610 - val\_loss: 0.6268  
Epoch 111/1000  
1500/1500 - 0s - loss: 0.4554 - val\_loss: 0.5956  
Epoch 112/1000

```

1500/1500 - 0s - loss: 0.4704 - val_loss: 0.5977
Epoch 113/1000
1500/1500 - 0s - loss: 0.4687 - val_loss: 0.5736
Epoch 114/1000
1500/1500 - 0s - loss: 0.4497 - val_loss: 0.5817
Epoch 115/1000
1500/1500 - 0s - loss: 0.4326 - val_loss: 0.5833
Epoch 116/1000
1500/1500 - 0s - loss: 0.4181 - val_loss: 0.5738
Epoch 117/1000
1500/1500 - 0s - loss: 0.4252 - val_loss: 0.5688
Epoch 118/1000
1500/1500 - 0s - loss: 0.4675 - val_loss: 0.5680
Epoch 119/1000
1500/1500 - 0s - loss: 0.4328 - val_loss: 0.5463
Epoch 120/1000
1500/1500 - 0s - loss: 0.4091 - val_loss: 0.5912
Epoch 121/1000
1500/1500 - 0s - loss: 0.4047 - val_loss: 0.5459
Epoch 122/1000
1500/1500 - 0s - loss: 0.4456 - val_loss: 0.5509
Epoch 123/1000
1500/1500 - 0s - loss: 0.4081 - val_loss: 0.5540
Epoch 124/1000
Restoring model weights from the end of the best epoch.
1500/1500 - 0s - loss: 0.4353 - val_loss: 0.5538
Epoch 00124: early stopping

```

Out[3]: <tensorflow.python.keras.callbacks.History at 0x1a40e6b0d0>

## Mean Square Error

The mean square error (MSE) is the sum of the squared differences between the prediction ( $\hat{y}$ ) and the expected ( $y$ ). MSE values are not of a particular unit. If an MSE value has decreased for a model, that is good. However, beyond this, there is not much more you can determine. We seek to achieve low MSE values. The following equation demonstrates how to calculate MSE.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

The following code calculates the MSE on the predictions from the neural network.

```

In [4]: from sklearn import metrics

# Predict
pred = model.predict(x_test)

# Measure MSE error.
score = metrics.mean_squared_error(pred, y_test)
print("Final score (MSE): {}".format(score))

```



Final score (MSE): 0.5463447829677607

## Root Mean Square Error

The root mean square (RMSE) is essentially the square root of the MSE. Because of this, the RMSE error is in the same units as the training data outcome. We desire Low RMSE values. The following equation calculates RMSE.

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$

```
In [5]: import numpy as np

# Measure RMSE error. RMSE is common for regression.
score = np.sqrt(metrics.mean_squared_error(pred,y_test))
print("Final score (RMSE): {}".format(score))
```

Final score (RMSE): 0.7391513938076291

## Lift Chart

We often visualize the results of regression with a lift chart. To generate a lift chart, perform the following activities:


- Sort the data by expected output and plot these values.
- For every point on the x-axis, plot that same data point's predicted value in another color.
- The x-axis is just 0 to 100% of the dataset. The expected always starts low and ends high.
- The y-axis is ranged according to the values predicted.

You can interpret the lift chart as follows:

- The expected and predict lines should be close. Notice where one is above the other.
- The below chart is the most accurate for lower ages.

```
In [7]: # Regression chart.
def chart_regression(pred, y, sort=True):
    t = pd.DataFrame({'pred': pred, 'y': y.flatten()})
    if sort:
        t.sort_values(by=['y'], inplace=True)
    plt.plot(t['y'].tolist(), label='expected')
    plt.plot(t['pred'].tolist(), label='prediction')
    plt.ylabel('output')
    plt.legend()
    plt.show()
```

```
# Plot the chart  
chart_regression(pred.flatten(),y_test)
```



In [ ]: